

Complete Architectural Planning Log: Project Millipede

This document preserves the entire chronological conversation and decision-making flow that established the architecture for **Project Millipede**.

Conversation Segment 1: The Monorepo Idea & Project Naming

User Prompt:

I have this idea. Kind of practicing everything that I'm learning and I love you in terms of software architecture back and long time testing is early everything into one big project. And so I'm not sure what to call the project basically called folder in which I'm going to do all of this but let's say I'm interested in the last and I will have a particular Antoine product being built and build it from the ground up from the backhand from the design the system architecture. And so while I'm designing and learning about it. But that having a component in AI or Neverland that I built from scratch. Because I'm learning about it. That should be that folder. Doing the client. Which is basically I want to play without me, you know learning and learning Kafka. I want to do kubernetes and daughter. I'm running about security. So I think I want to bring this whole experience that I'm learning. Which will make you go make me go deeper into each one of these topics or or things. And while I'm going to present learning, I'm actually Building Product because I I don't want to learn the vacuum. So I want to create bigger thing in the room all the things and to end. So what do I call the folder? How do I think about this right bring, you know the agents and ideas about you know, what I want to build from what I see at work or what I've seen out there. I want to have all of this in one particular space or folder or something help me here.

Response:

This is the ultimate developer rite of passage. Building a "God Project"—where you tie everything from Docker to custom AI models into one living, breathing ecosystem—is hands-down the best way to actually retain what you're learning. Learning in a vacuum is boring; building a playground where you can intentionally break and fix things is where the magic happens.

What to Name the Folder


The name should remind you of what this is: a grand experiment, a testing ground, and a production-grade product all at once. Here are a few vibes to choose from:

- **monorepo** or **workspace** : Simple, industry-standard. It tells your brain, "This is an enterprise-grade setup."
- **the-lab** or **sandbox** : Low pressure. It reminds you that it's okay if things break because it's a place for experimentation.
- **nexus** : The central link or connection. Perfect since you're tying Kafka, AI, and clients together.
- **foundry** or **the-crucible** : Where raw materials are melted down and forged into something strong.
- **singularity** : Because you are pulling *everything* into one single point.
- **project-hydra** : Cut off one bug, two more take its place.
- **playground-of-madness** : Honest and accurate for when Kubernetes starts acting up.

🧠 How to Think About This (The Mental Model)

To keep this from becoming a messy graveyard of half-finished tutorials, you need to treat this folder like a **Polyglot Monorepo**.

Instead of seeing it as one giant, tangled ball of yarn, structure it like an actual tech company would structure their codebase:



🚀 Reality Check: How to Not Burn Out

Because you want to build *everything* (AI from scratch, Kafka, Kubernetes, Security, etc.), the biggest risk here is **scope creep**. If you try to set up a Kubernetes cluster with a Kafka event stream and a custom neural network on day one, you will want to throw your laptop out the window.

Try a **layered approach**:

1. **Phase 1 (The Skeleton)**: Build a dead-simple backend and frontend. Get them talking.
2. **Phase 2 (The Container)**: Put them in Docker. Learn how Docker network works.
3. **Phase 3 (The Brains)**: Add your custom AI component. Let the backend talk to it.
4. **Phase 4 (The Scale)**: Introduce Kafka when you need to handle background jobs or data streaming asynchronously.
5. **Phase 5 (The Cloud/Orchestration)**: Move the whole thing into Kubernetes.

💬 Conversation Segment 2: Domain Ideation (Management, Cyber, Investing)

User Prompt:

i like olab for olsi lab (olsi being my name) or maybe millipede? (assuming im spelling it correctly). yes i like the oayered approach. I am in the cybersecurity space but i also like investing. i also like the idea of a e-commerce place as well not sure. what would be fun? i also want to build smth that scales for my job - i am a manager at work for example - i mange a team of 5 devs - anyythign that would help me build smth for management - that way it can be actally a real project at work to help other manages with stuff. anything - any suggestion is good actually. trying to brainstorm now

Response:

First off, both name choices are excellent. `olab` is clean and personal, while `millipede` is brilliant for a microservices project. A millipede has dozens of tiny legs (services) all moving independently, but working together to move the entire organism.

Here are three high-impact product concepts weaving together your management, security, and investing domains:

💡 Concept 1: "The DevSecOps Team Radar" (Management + Cyber)

An automated engineering management dashboard that monitors your team's development velocity *and* balances it against security risk. It hooks into GitHub, Jira, and security scanners, streaming commits, PRs, Jira transitions, and dependency alerts through Kafka in real time to provide team health insights without micromanagement.

💡 Concept 2: "FinOps Guarddog" (Management + Investing + Cyber)

In engineering, "investing" is how you allocate cloud budgets. This system monitors your infrastructure resource usage, cloud costs, and security posture on the Kubernetes cluster. An anomaly detection engine flags spikes where microservices consume abnormal bandwidth, indicating DDoS attacks or memory leaks.

💡 Concept 3: "The Intelligent Incident Commander" (Management + E-Commerce + Cyber)

A simulated high-traffic e-commerce store backed by an automated Incident Management system. When the store encounters database locks or SQL injection attempts, the system monitors logs via Kafka, diagnoses the issue, and drafts a Post-Mortem incident report for the manager.

💬 Conversation Segment 3: Deep Dive into DevSecOps Team Radar

User Prompt:

i like the 💡 Concept 1: "The DevSecOps Team Radar" (Management + Cyber) but I also like to build my own llm for scratch as im learning on manning.com. i also like wasm and solidJS and web components. i like to learn more about postgres and redis and other modern tools and frameworks in the modern microservices world. ill will give you more data to digest later.

Response:

Merging the DevSecOps Team Radar with a custom-built LLM, SolidJS, WebAssembly (Wasm), Postgres, and Redis elevates this to a highly sophisticated system.

The Stack Alignment:

1. **Frontend:** SolidJS + Web Components. SolidJS provides blazing-fast reactivity without a virtual DOM. Wasm compiled from Rust performs heavy local data parsing in the browser without UI freezes.
2. **Data & Streaming:** Kafka acts as the central event bus. Postgres is the long-term relational source of truth, and Redis manages real-time socket connections and ephemeral cache states.
3. **The Brains:** Your custom-built LLM. You wrap this PyTorch model in an API to parse semantic data from commit messages or reviews, analyzing sentiments, team friction, and burnout signs anonymously.

Conversation Segment 4: Choosing Backend Languages

User Prompt:

I want to use Rust. can we use more than one?

Response:

Absolutely. Using more than one language is the entire point of a microservices architecture! This is a **Polyglot Architecture**.

- **Where to Use Rust:** The high-throughput components. The Ingestion API (using Axum or Actix-web) and the Core Analyzer (consuming Kafka and interacting with Postgres/Redis) benefit immensely from Rust's performance and safety.
- **Where to Use Python:** The AI service. Since you are building an LLM from scratch, PyTorch and Python are perfect for managing mathematical tensor operations.

Conversation Segment 5: Security Posture

User Prompt:

what would one give us vs the other here: hould we look at building a custom security layer in Rust, or leverage something like JWTs passed through an API gateway?

Response:

- **JWTs via API Gateway:** Centralized edge bouncer. Offloads signature verification from internal services. However, it leads to a "hard shell, soft underbelly" layout where compromised internal nodes can spoof requests.
- **Custom Security Layer in Rust:** Decentralized Zero-Trust. Every node verifies the incoming cryptographic token. This handles Attribute-Based Access Control (ABAC) beautifully but increases code maintenance.
- **Hybrid Recommendation:** Build a custom, lightweight API Gateway in Rust. It decodes the public JWTs, signs an internal token, and routes requests to downstream nodes using internal Mutual TLS (mTLS).

Conversation Segment 6: mutual TLS (mTLS)

User Prompt:

i like the best of both worlds - we just ended up shipping mtls support at work.

Response:

Shipping mTLS in enterprise is huge. Implementing it at the code level in Rust (`rustls`) and Python (`ssl` / `pyOpenSSL`) inside your local monorepo provides deep knowledge.

- You generate a local Root CA via an OpenSSL script in `/infra` .
- Inject the signed certs and keys into Docker Compose.
- Build direct cryptographic handshakes between your custom Rust Gateway and your internal services.

Conversation Segment 7: Data Enrichment Pipeline

User Prompt:

the ai layer can intercept to enrich the data - wdyt?

Response:

This is the **Chained Consumer** or **Streaming ETL** pattern.

1. `raw-dev-events` topic receives webhooks from the Rust Ingestion API.
2. The custom Python LLM service consumes this, processes the text, adds sentiment and risk scores, and publishes to an `enriched-dev-events` topic.
3. The Rust Core Analyzer consumes the enriched topic, executing a single, clean database write to Postgres and updating the Redis cache.
4. **Benefit:** Log replayability. If you modify your LLM code, you can replay the immutable raw Kafka topic from offset zero and reconstruct your analytics without migrations.

Conversation Segment 8: The Data Inputs

User Prompt:

all of them plus maybe slack messages and logs from meetings? 1:1s - can that be?

Response:

This shifts the app from a simple dashboard to a **Developer Experience (DevEx) Context Engine**.

- **Git Diffs:** Analyzed by your LLM to identify high code churn.
- **Jira & PR Comments:** Analyzed to find communication bottlenecks or friction.
- **Slack & 1:1 Logs:** Sentiment analysis to identify burnout risks.
- **Security & Privacy:** Because 1:1 notes are highly personal, you must anonymize identifiers (synthetic pseudonymization) and store notes using field-level encryption.

Conversation Segment 9: Pulse Checks & Ingestion

User Prompt:

not kicking off anything yet . i want us to talk more about the plan. i have an idea about a little ping service - a slack ping to each of the team member like good morning and good afternoon at the end of the day to see the check in and checkout events and measure the pulse. wdyt? worth it? also how would i feed the events from slack - the 1:1s? since 1:1s data will live in google drive or maybe each one on one records etc

Response:

The Slack Pulse Check is worth it if framed as a "Friction Check" rather than an active micromanagement tool. It provides semi-structured sentiment data for the LLM.

- **Slack Ingestion:** Real-time push via Slack Events API webhooks to your Rust API Gateway.
- **1:1 Note Ingestion Options:**

- *Option A (Automated Scraper)*: Complicated background Cron job querying Google Drive API.
- *Option B (Manager-in-the-Loop)*: Paste notes into your dashboard. Secure, maintains a "human firewall," and allows you to strip highly private details manually before submission.

Conversation Segment 10: Client-Side WebAssembly Redaction

User Prompt:

i like option B

Response:

Option B is highly secure. You can build a local "Sanitize-on-Submit" component in your SolidJS frontend.

- The markdown editor offloads pasted text to a background browser Web Worker.
- The Web Worker runs a Rust WebAssembly binary compiled with `wasm-pack`.
- The Wasm engine tokenizes names, emails, and sensitive keys, returning match offsets.
- The UI displays interactive highlights, letting you approve or adjust redact markers.
- Deterministic pseudonymization replaces "Alice" with `Developer_Alpha` before the data leaves the boundary.

Conversation Segment 11: SolidJS vs. React

User Prompt:

another question - are we sure about solidJS vs React here? i like solidJS but anything makes us reconsider? i think Python AI microservice runs asynchronously in the background as a continuous Kafka consumer ? wdyt?

Response:

- **SolidJS Wins:** Compilation to raw DOM nodes (no Virtual DOM) matches beautifully with WebAssembly buffer sharing and Web Component isolation.
- **React Considerations:** React's massive dashboard and chart library ecosystem (`Tremor` , `Recharts` , `shadcn/ui`) speeds up complex visualizations. Additionally, using React aligns with your team's stack at work.
- **SolidJS Decision:** Since you are prioritizing learning, SolidJS is highly rewarding. You can bind vanilla JS chart libraries (like `ApexCharts`) directly inside Solid's `onMount` primitive.

Conversation Segment 12: System Roadmap

The system implementation maps to a cohesive 4-phase master plan:

- **Phase 1 (Skeleton):** Run Kafka, Postgres, and Redis inside Docker Compose. Construct an Axum-based Rust Ingestion endpoint that publishes raw events to Kafka.
- **Phase 2 (Security):** Set up the custom Rust API Gateway running `rustls` to route traffic over internal mTLS connections.
- **Phase 3 (AI Integration):** Build the Python AI microservice consuming from Kafka, executing inference via PyTorch, and piping results to the enriched stream.

- **Phase 4 (Frontend):** Build the SolidJS dashboard featuring the Web Worker Rust-Wasm redaction module.